

BUSINESS OVERVIEW

TR-Insight

AI QE that knows your results — how it solves the test-reporting problem, for the people who fund it and rely on it.

Product: TR-Insight (part of the TR-TestIntelligence suite)

Publisher: Testrig Technologies

Audience: Executives, QA Leadership, Test Managers, Procurement

Document version: 1.0 | **Date:** September 2, 2026

Table of Contents

1	Executive Summary
2	The Problem TR-Insight Solves
3	How TR-Insight Solves It
4	Who It's For
5	What You Can Do With It
6	Bring the Data You Already Have
7	Why AI, and Why It's Safe to Rely On
8	Cost You Can Predict
9	Data Security & Ownership
10	Getting Started — What It Takes
11	Deployment Choices & Indicative Cost
12	Roles & Access, in Plain Terms
13	The Outcome
14	Part of a Larger Suite

1. Executive Summary

Your team already runs tests and already produces results. What's missing isn't more testing — it's a fast, trustworthy way to turn those results into an answer a leader can act on. **TR-Insight** is an AI-powered test analytics dashboard that sits on top of the test results you already have and answers plain-English questions about them — instantly, with a chart, and with a release call you can defend.

Minutes

not afternoons, to answer a release-
readiness question

0

manual spreadsheet exports required

4

audiences served from one source of
truth

It runs entirely inside your own environment — nothing about your releases, defects, or customer-sensitive test data leaves your infrastructure. It is deployed with a single command and starts answering questions the same day.

2. The Problem TR-Insight Solves

Every QA organization running at any real scale eventually runs into the same five problems:

- **No defensible answer, under time pressure.** A release is hours away and nobody can say, with evidence, whether it's safe to ship.
- **Reporting is manual labor.** Every status update starts with someone exporting results into a spreadsheet by hand.
- **One result, two stories.** The same failure means one thing to leadership and something else entirely to the engineer who owns the suite — there's no shared, role-appropriate framing.
- **Flaky tests hide in the noise** until they take down a release nobody saw coming.
- **Tribal knowledge is a single point of failure.** Insight depends on the one person who knows where the data lives — and they're on leave the week it matters.

"We can't tell the CTO whether we're safe to ship without someone spending an hour building a slide first." — the exact situation TR-Insight is built to eliminate.

3. How TR-Insight Solves It

Each problem above maps directly to a capability in the product:

PROBLEM

No defensible go/no-go under time pressure

TR-INSIGHT

A clear release-readiness verdict, with the failing evidence one click away when someone challenges it

PROBLEM

Manual spreadsheet exports for every status update

TR-INSIGHT

Summaries and charts produced on demand, in plain English, with no reporting chore

PROBLEM

Same failure, different story for different audiences

TR-INSIGHT

One set of results, role-aware answers — executive risk framing for leadership, full engineering detail for the team that owns the suite

PROBLEM

Flaky tests hiding in the noise

TR-INSIGHT

Regressions, flaky tests and recurring patterns surfaced automatically across builds, ranked by impact

PROBLEM

Insight locked inside one person's head

TR-INSIGHT

Anyone with a login can self-serve — analysis stops being a bottleneck tied to one individual

4. Who It's For

TR-Insight is built around the questions different people actually open a dashboard to answer — not a one-size-fits-all report.

CTO / Engineering Leadership

"Are we safe to ship?"

Release risk and quality trend in plain language — a defensible go/no-go in one screen, no deck required.

QA Lead

"Where is quality actually weak?"

The modules dragging the pass rate down, ranked by impact — effort goes where it actually changes the number.

Test Manager

"What do I report this week?"

Coverage, pass rate and regressions across builds, on demand — reporting time collapses to a question, asked and answered.

SDET / QA Engineer

"Why did this fail?"

The failing tests, their errors, the flaky ones and the slowest ones — straight to what's actually worth debugging.

5. What You Can Do With It

- **Ask in plain English** — e.g. "is the payment module stable?" — and get a direct, sourced answer, not a raw data dump.
- **Get charts and dashboards on demand** — presentation-ready visuals, ready to drop into a status update or a steering call, no design effort required.
- **See trends and patterns automatically** — recurring failures, regressions and flaky tests found, explained, and ranked, with a recommendation on what to fix first.
- **Get answers shaped for the reader** — leadership gets risk and release framing, engineers get failing tests and technical detail, all from the same underlying source of truth.
- **Defend a go/no-go call** — a clear release-readiness verdict, with the evidence one click away if someone pushes back.
- **Watch a test run live** — follow execution as it happens, so problems surface during the run itself instead of in tomorrow's report.

6. Bring the Data You Already Have

There is no migration project and no rebuild required. TR-Insight connects, in a few clicks, to what your team is already producing:

- Automation reports (from your existing test frameworks)
- Spreadsheet exports and Excel workbooks maintained by hand
- Existing databases
- Internal reporting services, via API
- Live test runs, streamed while they're still executing

Four practical routes get data in — a watched drop-box folder, a direct upload endpoint for CI systems that can't share a filesystem, a manual "Add Build" action inside the dashboard, or a one-shot bulk import. Whichever route is used, the result is the same: a new build appears in the dashboard, typically within seconds, with no restart or manual refresh needed.

7. Why AI, and Why It's Safe to Rely On

Why AI belongs here

- Test data is large, repetitive, and dull to read line-by-line — exactly the work people postpone, and exactly what AI does instantly.
- It reads *every* result, not the sample someone had time to look at, so nothing quietly gets skipped.
- It explains what happened in language a non-technical stakeholder can act on, with no translation layer required.
- It ranks what actually matters — the one failure blocking the release, ahead of the twenty that don't.

Why the answers can be trusted

Charts and figures are *not* free-form AI output. The system deterministically parses what chart is being asked for, has the AI write only the underlying data query, and then validates and downgrades the result if the data can't actually support the requested chart form. In plain terms: the AI decides *what to look up*, but a fixed, testable process decides *what gets drawn* — so the same question reliably produces the same kind of answer, and visuals can't silently misrepresent the underlying numbers.

8. Cost You Can Predict

- **Only what's needed reaches the AI** — never the entire dataset for every question, which keeps per-question cost small and consistent.
- **Repeat and follow-up questions reuse work already done** instead of paying for the same computation twice.
- **Every user and team has a spend limit**, so AI cost cannot run away unnoticed between invoices.
- **Usage is visible per person and per project** from an admin console, so what AI costs — and why — is never a surprise at the end of the month.

AI spend becomes a known, budgeted line item — not an open-ended bill that leadership discovers after the fact.

9. Data Security & Ownership

TR-Insight runs **inside your own environment** — on your own server, VM, or cloud account, brought up with a single command. Release information, defect detail, and customer-sensitive test data never leave your infrastructure by default, which keeps security review short and procurement simple compared to a SaaS product that requires sending your data to a third party for storage.

- Accounts are protected with industry-standard authentication and password hashing; every deployment gets its own unique credentials — there is no shared factory password across installs.
- Access is scoped by **team** ("workspace") — a person only sees the builds that belong to their own team unless they hold an administrator or leadership role.
- Every significant action — logins, account changes, permission changes, deletions — is recorded in a permanent audit trail that leadership and compliance can review at any time.
- The one place data does travel is to the AI provider your organization chooses (e.g., Gemini, OpenAI, Anthropic, or a fully local/offline model via Ollama) — and only the specific question and the minimal supporting data it needs, never the full dataset.

10. Getting Started — What It Takes

Standing TR-Insight up does not require a services engagement or a lengthy implementation project:

1. One setup command brings up the entire application (dashboard + backend) on a server, VM, or laptop.
2. The first login is a one-time administrator account, which is immediately forced to set a real password.
3. The administrator picks an AI provider from a short list (or connects a fully private, offline model) — this can also be finished later without restarting anything.
4. The first set of test results is dropped in, and it appears in the dashboard within about 30 seconds.
5. Everyone else requests an account through a simple self-service sign-up, approved by an administrator into the right team.

From a cold start to a working, populated dashboard is typically a same-day effort for a technical team member — not a multi-week rollout.

11. Deployment Choices & Indicative Cost

Option	Best for	Indicative running cost
Single server / VM you already own	Teams with existing infrastructure and IT support	No extra cloud spend
A single cloud server (e.g., AWS), stood up via an included automated setup	Teams wanting a supported, documented, self-hosted path with minimal setup effort	Roughly \$30–\$60/month for the server, plus a few dollars/month for storage
Split hosting — dashboard on a managed hosting service, backend on your own server	Teams who want the dashboard itself managed and updated automatically	Backend hosting cost only; the managed frontend hosting has its own separate (often free-tier-eligible) pricing

In every option, all of your test data and accounts live on infrastructure your organization controls — the deployment choice only affects convenience and operating cost, never who owns the data.

12. Roles & Access, in Plain Terms

Role	What they can do
Administrator / CTO	See everything across every team, manage user accounts, configure the AI provider, view the full audit history
QA Manager	Full read and import access, scoped to their own team's results
SDET / QA Engineer	Full read and import access, scoped to their own team's results

A "team" (called a workspace) is simply a boundary around which builds a person can see — assigning someone to a new team is all it takes to create one; there's no separate setup step. New sign-ups sit in a pending queue until an administrator approves them into the correct team and role, so access is always a deliberate decision, never automatic.

Independently of account access, each person can also choose which "voice" the AI answers in — an executive risk summary versus full technical failure detail — without that choice changing what data they're allowed to see.

13. The Outcome

- **Hours back every week** — the export-and-summarize cycle disappears; questions that used to take an afternoon are answered while the meeting is still running.
- **Decisions you can defend** — every answer carries the data it came from, so release calls stop being a matter of opinion.
- **Problems caught earlier** — regressions and unstable areas surface while a fix is still cheap, not after they've already cost a release.
- **AI cost you can predict** — spend is capped per person and per team, a known line item rather than an open-ended bill.

14. Part of a Larger Suite

TR-Insight is one piece of Testrig Technologies' **TR-TestIntelligence** suite of AI-driven QA tools, designed to work together across the testing lifecycle:

Product	What it does
TR-Forge	AI-assisted generation of test cases
TR-Automate	AI-assisted Playwright test automation
TR-Insight	Test reporting and analytics (this document)

Together, the suite covers writing tests, running them, and understanding the results — with AI applied at each stage where it removes manual, repetitive work rather than adding process.